

Lab 4: Healthy Momma, Healthy Baby

Krista Mar and Nikki Haas

12/1/2016

A Nice Introduction that Makes Us Sound Like Pros

[illegible]

Step 1: Read in the Data

```
load('/Users/nicholeh/student285/w203/w203_lab_4/bwght_w203.RData')
desc
```

##	variable	label
## 1	mage	mother's age, years
## 2	meduc	mother's educ, years
## 3	monpre	month prenatal care began
## 4	npvis	total number of prenatal visits
## 5	fage	father's age, years
## 6	feduc	father's educ, years
## 7	bwght	birth weight, grams
## 8	omaps	one minute apgar score
## 9	fmaps	five minute apgar score
## 10	cigs	avg cigarettes per day
## 11	drink	avg drinks per week
## 12	lbw	=1 if bwght <= 2000
## 13	vlbw	=1 if bwght <= 1500
## 14	male	=1 if baby male
## 15	mwhte	=1 if mother white
## 16	mbldk	=1 if mother black
## 17	moth	=1 if mother is other
## 18	fwhte	=1 if father white
## 19	fbldk	=1 if father black
## 20	foth	=1 if father is other
## 21	lbwght	log(bwght)
## 22	agesq	mage ²
## 23	npvissq	npvis ²

Step 2: Exploratory Data Analysis

First, get summary statistics on each element of the dataset:

```
nrow(data)
```

```
## [1] 1832
```

```
summary(data)
```

```
##      mage      meduc      monpre      npvis
## Min.   :16.00   Min.    : 3.00   Min.    :0.000   Min.    : 0.00
## 1st Qu.:26.00   1st Qu.:12.00   1st Qu.:1.000   1st Qu.:10.00
## Median :29.00   Median :13.00   Median :2.000   Median :12.00
## Mean   :29.56   Mean    :13.72   Mean    :2.122   Mean    :11.62
## 3rd Qu.:33.00   3rd Qu.:16.00   3rd Qu.:2.000   3rd Qu.:13.00
## Max.   :44.00   Max.    :17.00   Max.    :9.000   Max.    :40.00
##      NA's      :30      NA's    :5      NA's    :68
##      fage      feduc      bwght      omaps
## Min.   :18.00   Min.    : 3.00   Min.    : 360   Min.    : 0.000
## 1st Qu.:28.00   1st Qu.:12.00   1st Qu.:3076   1st Qu.: 8.000
## Median :31.00   Median :14.00   Median :3425   Median : 9.000
## Mean   :31.92   Mean    :13.92   Mean    :3401   Mean    : 8.386
## 3rd Qu.:35.00   3rd Qu.:16.00   3rd Qu.:3770   3rd Qu.: 9.000
## Max.   :64.00   Max.    :17.00   Max.    :5204   Max.    :10.000
##      NA's      :6      NA's    :47      NA's    :3
##      fmaps      cigs      drink      lbw
## Min.   : 2.000   Min.    : 0.000   Min.    :0.0000   Min.    :0.00000
## 1st Qu.: 9.000   1st Qu.: 0.000   1st Qu.:0.0000   1st Qu.:0.00000
## Median : 9.000   Median : 0.000   Median :0.0000   Median :0.00000
## Mean   : 9.004   Mean    : 1.089   Mean    :0.0198   Mean    :0.01638
## 3rd Qu.: 9.000   3rd Qu.: 0.000   3rd Qu.:0.0000   3rd Qu.:0.00000
## Max.   :10.000   Max.    :40.000   Max.    :8.0000   Max.    :1.00000
##      NA's      :3      NA's    :110   NA's    :115
##      vlbw      male      mwhte      mblick
## Min.   :0.000000   Min.    :0.0000   Min.    :0.0000   Min.    :0.0000
## 1st Qu.:0.000000   1st Qu.:0.0000   1st Qu.:1.0000   1st Qu.:0.0000
## Median :0.000000   Median :1.0000   Median :1.0000   Median :0.0000
## Mean   :0.007096   Mean    :0.5136   Mean    :0.8865   Mean    :0.0595
## 3rd Qu.:0.000000   3rd Qu.:1.0000   3rd Qu.:1.0000   3rd Qu.:0.0000
## Max.   :1.000000   Max.    :1.0000   Max.    :1.0000   Max.    :1.0000
##
##      moth      fwhite      fblack      foth
## Min.   :0.00000   Min.    :0.0000   Min.    :0.0000   Min.    :0.0000
## 1st Qu.:0.00000   1st Qu.:1.0000   1st Qu.:0.0000   1st Qu.:0.0000
## Median :0.00000   Median :1.0000   Median :0.0000   Median :0.0000
## Mean   :0.05404   Mean    :0.8897   Mean    :0.05841   Mean    :0.05186
## 3rd Qu.:0.00000   3rd Qu.:1.0000   3rd Qu.:0.0000   3rd Qu.:0.0000
## Max.   :1.00000   Max.    :1.0000   Max.    :1.0000   Max.    :1.0000
##
##      lbwght      magesq      npvissq
## Min.   :5.886   Min.    : 256.0   Min.    : 0.0
## 1st Qu.:8.031   1st Qu.: 676.0   1st Qu.:100.0
## Median :8.139   Median : 841.0   Median :144.0
## Mean   :8.114   Mean    : 896.4   Mean    :148.6
## 3rd Qu.:8.235   3rd Qu.:1089.0   3rd Qu.:169.0
## Max.   :8.557   Max.    :1936.0   Max.    :1600.0
##      NA's      :68
```

Response Variables

The bwght, lbwght, omaps and fmaps variables are related to the health of the baby.

The first thing to check is if these variables are collinear. We will omit bwghts as that is a function of lbwghts.

```
library(ggplot2)
cor(data$omaps, data$fmaps, use = "complete.obs")
```

```
## [1] 0.5575238
```

```
cor(data$lbwght, data$fmaps, use = "complete.obs")
```

```
## [1] 0.2710456
```

```
p <- ggplot(data, aes(omaps, lbwght)) + geom_point(size = 0.25) +
  geom_smooth(method = "lm", se = FALSE) + geom_point(aes(colour = fmaps)) +
  ggtitle("Scatterplot of log(weight) against One Minute APGAR test,\n
          with 5 minute APGAR test heatmap")
```


p

```
## Warning: Removed 3 rows containing non-finite values (stat_smooth).
```


```
## Warning: Removed 3 rows containing missing values (geom_point).
```

```
## Warning: Removed 3 rows containing missing values (geom_point).
```

Scatterplot of log(weight) against One Minute APGAR test,
with 5 minute APGAR test heatmap



```
p <- ggplot(data, aes(factor(fmaps), lbwght)) + geom_boxplot()
p
```



Look at the extreme fmops case

```
data[data$fmops < 4,]
```

```
##      mage meduc monpre npvis fage feduc bwght omaps fmops cigs drink lbw
## NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
## 837     32     12       2     10     40     16    2580      2      2      0      0      0
## NA.1    NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
## NA.2    NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
##      vlbw male mwhite mblack moth fwhte fblack foth   lbwght magesq npvissq
## NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
## 837      0      1      1      0      0      1      0      0 7.855545  1024     100
## NA.1    NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
## NA.2    NA      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
```

Looking at the data, we can be reasonably assured that the response variables are related, but not collinear. It may be best to make a combined variable of `fmops` and `omaps` such as `mapscombined = fmops + omops`. The difference would not make much sense compared to the sum; 10 - 10 and 2 - 2 are both zero, after all.

Regressors

The variables `monpre` and `npvis` are related to the prenatal care given during pregnancy. Let us review them for collinearity:

```
cor(data$npvis, data$monpre, use = "complete.obs")
```


```
## [1] -0.3061006
```

```
ggplot(data, aes(monpre, npvis)) + geom_point(size = 0.25) +  
  geom_smooth(method = "lm", se = FALSE) + geom_jitter()
```

```
## Warning: Removed 69 rows containing non-finite values (stat_smooth).
```


```
## Warning: Removed 69 rows containing missing values (geom_point).
```

```
## Warning: Removed 69 rows containing missing values (geom_point).
```



```
ggplot(data, aes(factor(monpre), npvis)) + geom_boxplot()
```

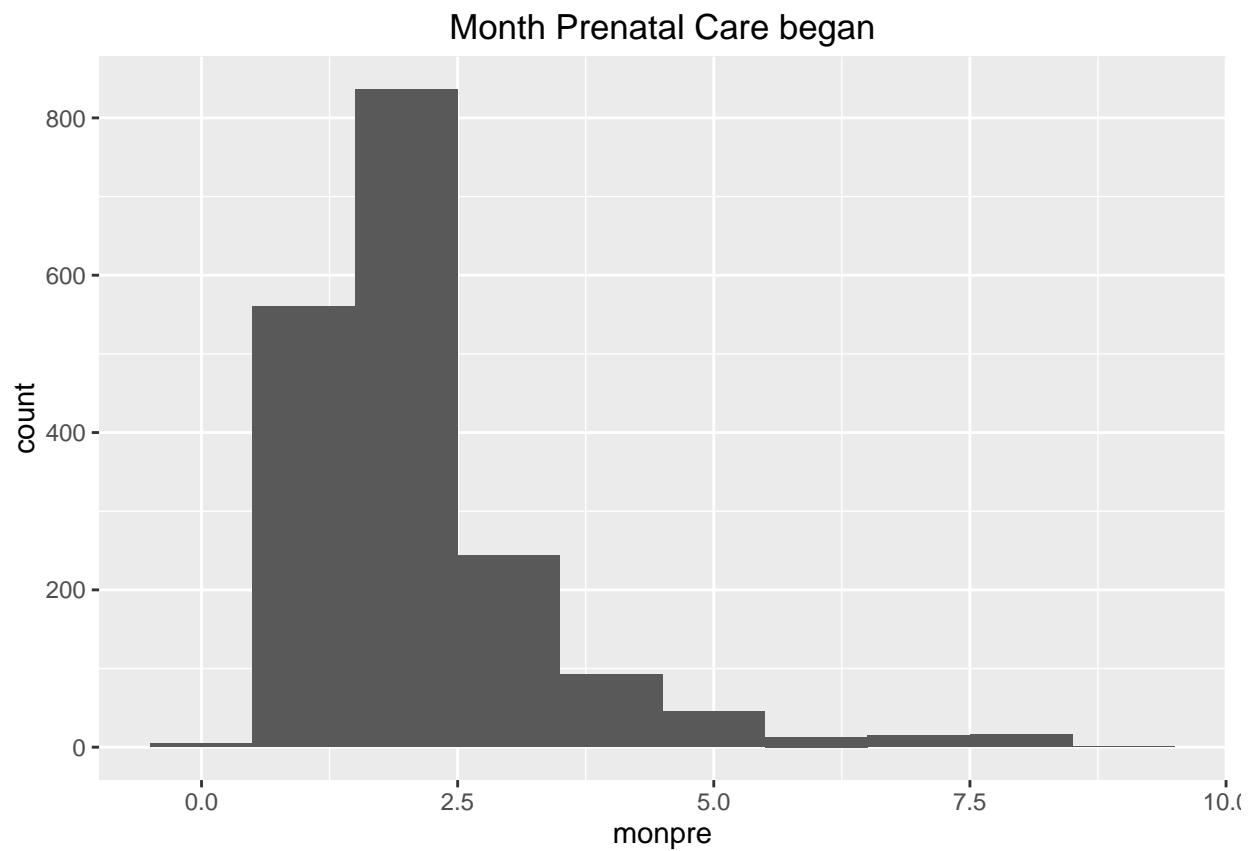
```
## Warning: Removed 68 rows containing non-finite values (stat_boxplot).
```



From this set, we can see that the data is not collinear, and indeed we can see that we might have some reporting errors. 5 mothers are listed as starting prenatal care in month 0 of their pregnancy, but they visited the doctor 0 times. These probably denote missing information or an error in reporting. Unfortunately, this data does show a definitive downward trend leading us to suspect that the number of visits is a function of month prenatal care began. This makes sense intuitively; if a mother starts prenatal care in her 2nd month of pregnancy, she has ample time for frequent doctor visits. However, if she starts her prenatal care towards the end of her pregnancy, she does not have enough time to visit the doctor as often as a woman who started in month 2.


```
ggplot(data, aes(x=monpre)) + geom_histogram(aes(y = ..count..),bins = 10) +
  ggtitle("Month Prenatal Care began")
```

```
## Warning: Removed 5 rows containing non-finite values (stat_bin).
```




```
ggplot(data, aes(x=sqrt(monpre))) + geom_histogram(aes(y = ..count..), bins = 10) +  
  ggtitle("Month Prenatal Care Began, Half Power")
```

```
## Warning: Removed 5 rows containing non-finite values (stat_bin).
```




```
ggplot(data, aes(x=log(monpre))) + geom_histogram(aes(y = ..count..), bins = 10) +  
  ggtitle("Month Prenatal Care Began, Natural Log")
```

```
## Warning: Removed 10 rows containing non-finite values (stat_bin).
```


```
ggplot(data, aes(x=(monpre^2))) + geom_histogram(aes(y = ..count..), bins = 10) +  
  ggtitle("Month Prenatal Care Began, Square Power")
```

```
## Warning: Removed 5 rows containing non-finite values (stat_bin).
```



```
ggplot(data, aes(x=npvis)) + geom_histogram(aes(y = ..count..), bins = 15) +  
  ggtitle("Number of Prenatal Visits")
```

```
## Warning: Removed 68 rows containing non-finite values (stat_bin).
```



All in all, the number of visits follows a mostly normal curve, and the square root of the month prenatal care began follow a mostly normal curve. Then we say smart things about how that will all relate to each other.

Step 3: Modeling

Model 1: Basic Linear Model

First we will check variants on the requested model; baby health as a function of the mother's prenatal care. We will not

```
model1 = lm(data$bwght ~ data$monpre + data$npvis)
model1

##
## Call:
## lm(formula = data$bwght ~ data$monpre + data$npvis)
##
## Coefficients:
## (Intercept)  data$monpre  data$npvis
##      3161.27      17.06      17.55

model2 = lm(data$lbwght ~ data$monpre + data$npvis)
model2

##
## Call:
## lm(formula = data$lbwght ~ data$monpre + data$npvis)
##
## Coefficients:
```

```
## (Intercept) data$monpre data$npvis
##      8.008629      0.008570      0.007503

model3 = lm(data$bwght ~ sqrt(data$monpre) + data$npvis)
model3

##
## Call:
## lm(formula = data$bwght ~ sqrt(data$monpre) + data$npvis)
##
## Coefficients:
##      (Intercept)      sqrt(data$monpre)      data$npvis
##          3122.10              54.93              17.39

model4 = lm(data$lbwght ~ sqrt(data$monpre) + data$npvis)
model4

##
## Call:
## lm(formula = data$lbwght ~ sqrt(data$monpre) + data$npvis)
##
## Coefficients:
##      (Intercept)      sqrt(data$monpre)      data$npvis
##          7.988854              0.027646              0.007423

AIC(model1)

## [1] 27428.8

AIC(model2)

## [1] -603.5921

AIC(model3)

## [1] 27428.8

AIC(model4)

## [1] -603.6182
```

Model 2: An Alternate Main Model

The 1 minute and 5 minute APGAR scores on their own do not tell us much. As we can see from the heatmap on the first scatterplot, a baby who has a low one minute score tends to have a higher five minute score. There are very few examples of a baby having a worse five minute score than a one minute score:

```
nrow(data[!is.na(data$fmaps) < !is.na(data$omaps),])
```

```
## [1] 3
```

However, we can get some information if we sum up `omaps` and `fmaps`. A baby that goes from 0 to 10 then would have an overall low score compared to a baby who started with a score of 10 and was still at 10 5 minutes later.

```
data$combined_apgarscores = data$omaps + data$fmaps
```

Now that we have a calculated field that sums up the APGAR tests, we can try an alternate linear model:

```
a1 = lm(data$combined_apgarscores ~ data$monpre + data$npvis)
a1
```

```
##
## Call:
## lm(formula = data$combined_apgarscores ~ data$monpre + data$npvis)
##
## Coefficients:
## (Intercept) data$monpre data$npvis
## 17.08460 -0.03600 0.03303
a2 = lm(data$combined_apgarscores ~ data$monpre)
a3 = lm(data$combined_apgarscores ~ data$npvis)
summary(a1)
```

```
##
## Call:
## lm(formula = data$combined_apgarscores ~ data$monpre + data$npvis)
##
## Residuals:
## Min 1Q Median 3Q Max
## -13.343 -0.409 0.555 0.624 2.792
##
## Coefficients:
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) 17.084596 0.149953 113.933 < 2e-16 ***
## data$monpre -0.036001 0.029285 -1.229 0.219117
## data$npvis 0.033032 0.009818 3.364 0.000783 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.441 on 1757 degrees of freedom
## (72 observations deleted due to missingness)
## Multiple R-squared: 0.009575, Adjusted R-squared: 0.008448
## F-statistic: 8.493 on 2 and 1757 DF, p-value: 0.0002134
```

```
AIC(a1)
```

```
## [1] 6285.572
```

```
AIC(a2)
```



```
## [1] 6518.486
```



```
AIC(a3)
```

```
## [1] 6288.508
```

Our AIC scores tell us that a1 is efficient than the other three, so let us explore a1 further:

```
plot(a1)
```






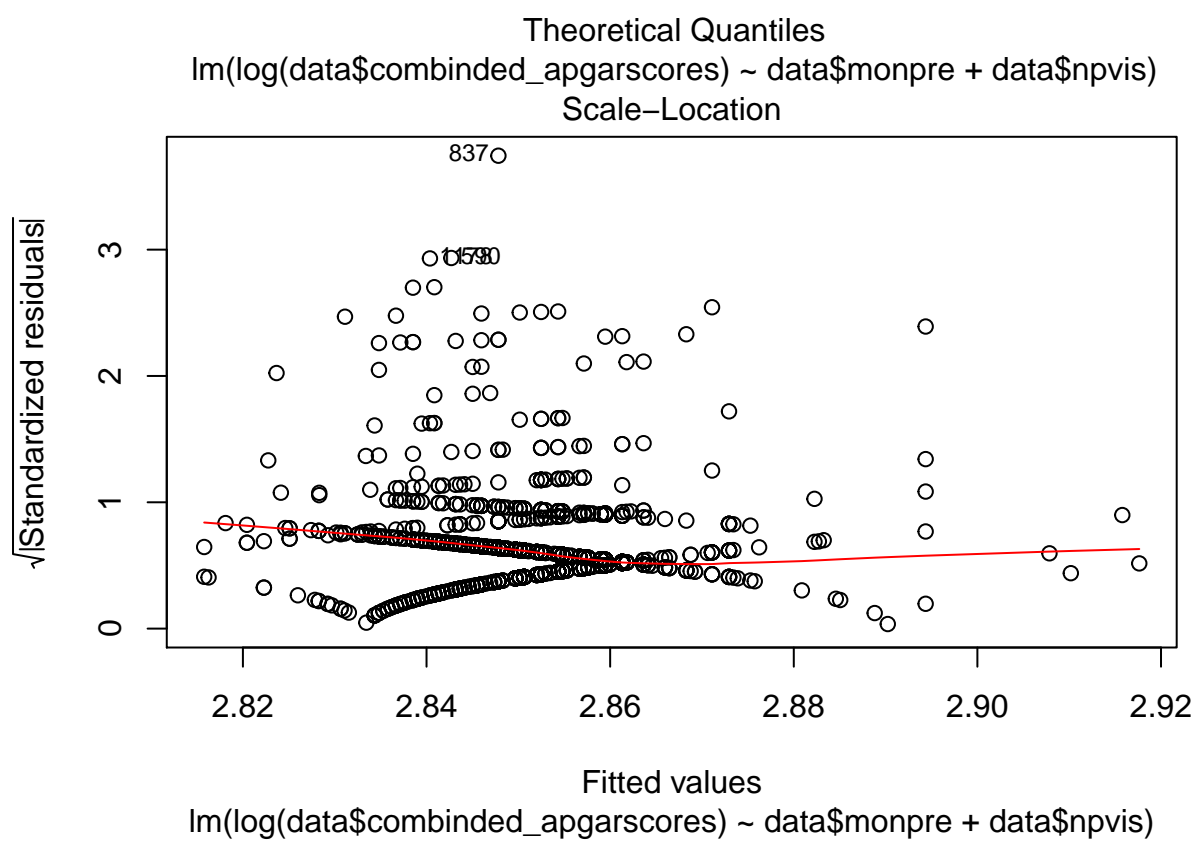

```
## lm(formula = log(data$combined_apgarscores) ~ data$monpre +
##     data$npvis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.46153 -0.01927  0.03509  0.04022  0.15674
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.8282616   0.0108567  260.510 < 2e-16 ***
## data$monpre  -0.0018509   0.0021202  -0.873  0.38280
## data$npvis    0.0023268   0.0007108   3.273  0.00108 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1043 on 1757 degrees of freedom
## (72 observations deleted due to missingness)
## Multiple R-squared:  0.008253, Adjusted R-squared:  0.007124
## F-statistic: 7.311 on 2 and 1757 DF, p-value: 0.0006889
```


```
AIC(a4)
```

```
## [1] -2956.35
```

```
plot(a4)
```







$\text{lm}(\log(\text{data\$combined_apgarscores}) \sim \text{data\$monpre} + \text{data\$npvis})$

Again, we didn't get very good results. We cannot look at the logs of monpre or npvis as they contain some zero values, and we suspect those datapoints are important as they appear to represent mothers who had no prenatal care. Instead, let us try the product of the APGAR scores.

```
data$product_apgarscores = data$omaps * data$fmaps
```



```
a5 = lm(data$product_apgarscores ~ data$monpre + data$npvis)
summary(a5)
```

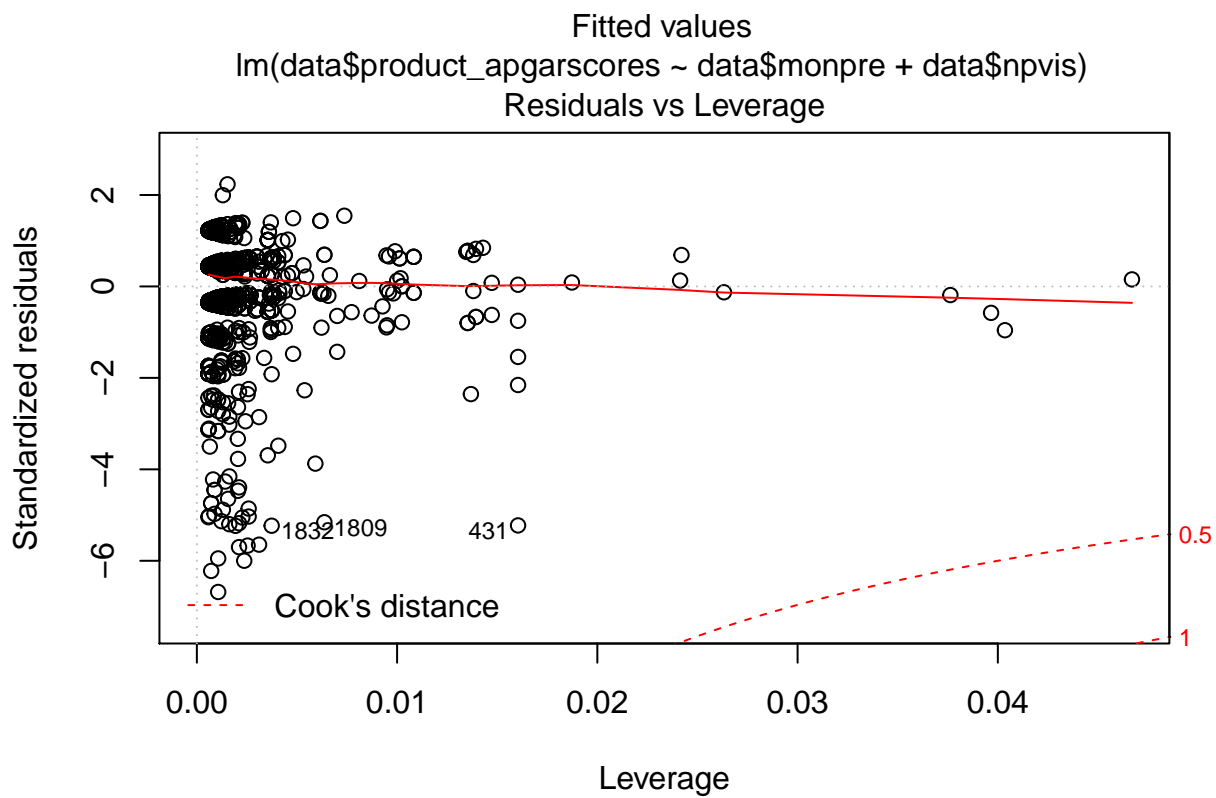

```
##
## Call:
## lm(formula = data$product_apgarscores ~ data$monpre + data$npvis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -76.741  -3.975   4.681   5.280  25.645
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  73.59915    1.19580  61.548 < 2e-16 ***
## data$monpre  -0.34388    0.23353  -1.472  0.14107
## data$npvis    0.25532    0.07829   3.261  0.00113 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 11.49 on 1757 degrees of freedom
## (72 observations deleted due to missingness)
## Multiple R-squared:  0.00981,    Adjusted R-squared:  0.008683
## F-statistic: 8.704 on 2 and 1757 DF,  p-value: 0.0001732
```

```
AIC(a5)
```

```
## [1] 13593.95
```

```
plot(a5)
```





This data is not looking good either. Let's try removing the non-normally distributed `data$monpre` field

```
a6 = lm(data$product_apgarscores ~ data$npvis)
summary(a6)
```


```
##
## Call:
```

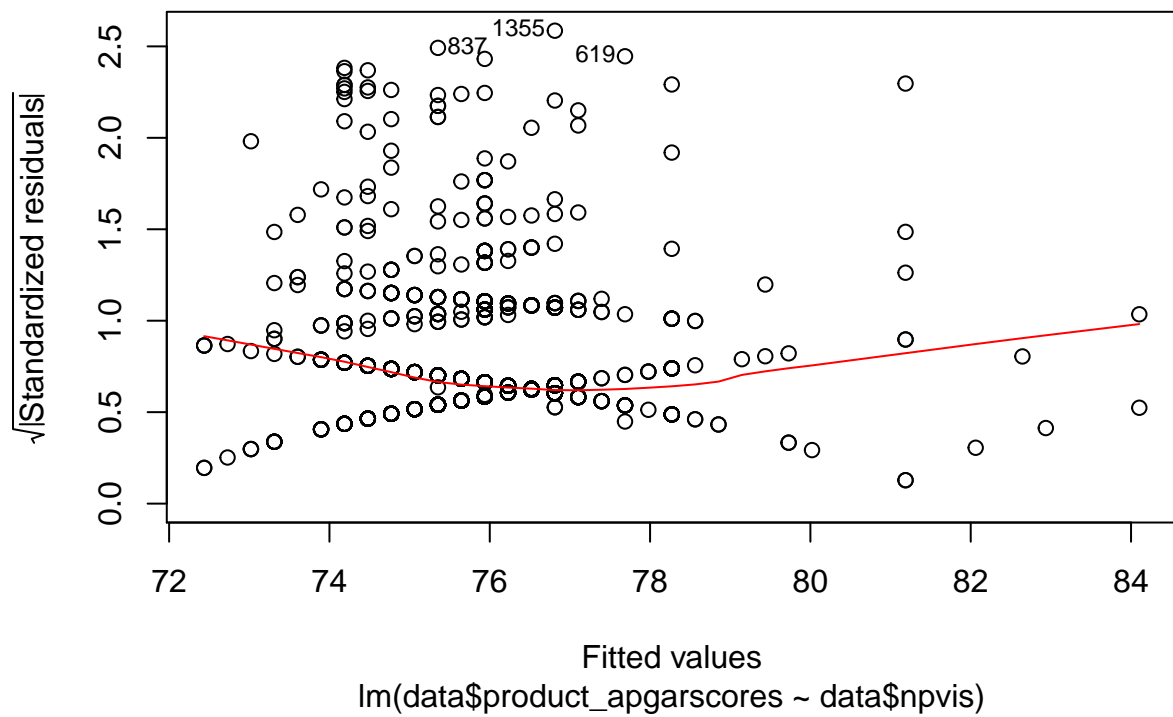
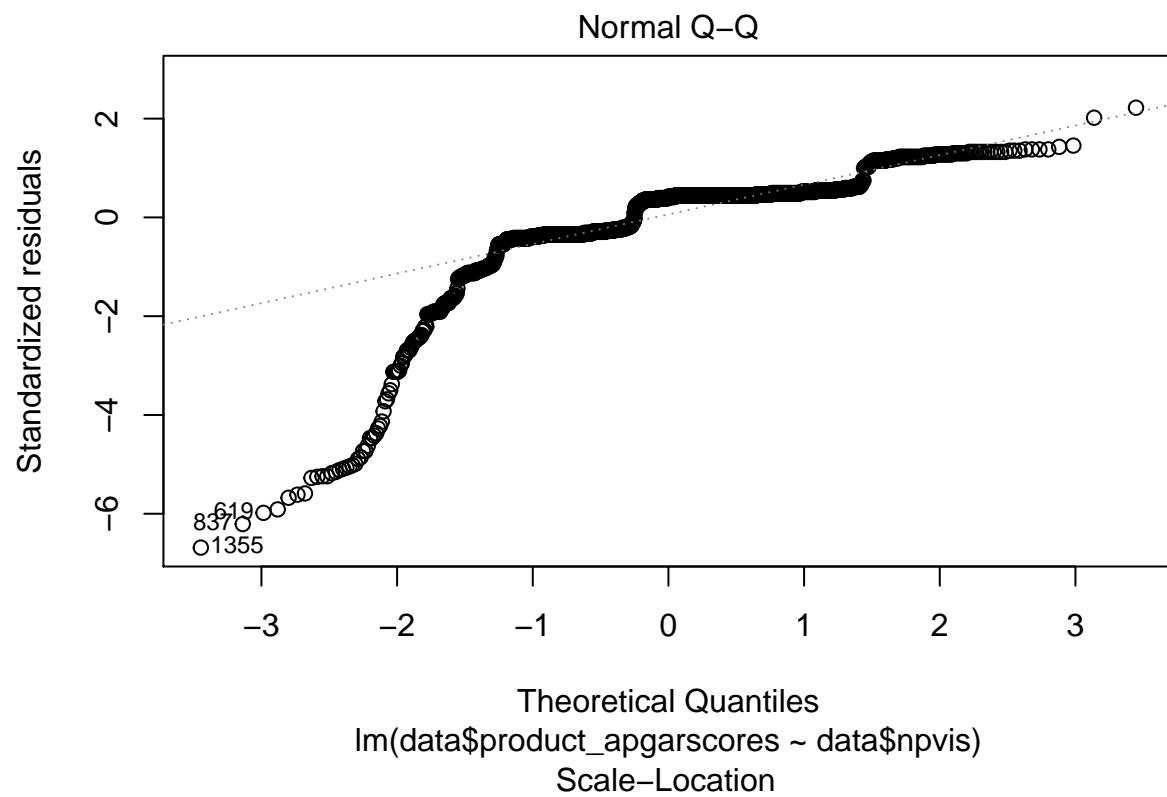
```
## lm(formula = data$product_apgarscores ~ data$npvis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -76.812  -3.937   4.771   5.354  25.521
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  72.43739    0.90871   79.715 < 2e-16 ***
## data$npvis    0.29165    0.07452    3.914 9.43e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 11.5 on 1759 degrees of freedom
## (71 observations deleted due to missingness)
## Multiple R-squared:  0.008633, Adjusted R-squared:  0.00807
## F-statistic: 15.32 on 1 and 1759 DF, p-value: 9.429e-05
```


```
AIC(a6)
```

```
## [1] 13601.99
```

```
plot(a6)
```







Unfortunately that too is a little worse. Let's try normalizing these calculated values:

```
data$normalized_combined_apgar = (data$combined_apgarscores - mean(!is.na(data$combined_apgarscores)))
data$normalized_product_apgar = (data$product_apgarscores - mean(!is.na(data$product_apgarscores)))/sd(
a7 = lm(data$normalized_combined_apgar~data$monpre + data$npvis)
a8 =lm(data$normalized_product_apgar~data$monpre + data$npvis)
summary(a7)
```

```
##
## Call:
## lm(formula = data$normalized_combined_apgar ~ data$monpre + data$npvis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -329.91  -10.11   13.72   15.43   69.04
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  397.7350     3.7076 107.275 < 2e-16 ***
## data$monpre  -0.8901     0.7241  -1.229 0.219117
## data$npvis    0.8167     0.2428   3.364 0.000783 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 35.63 on 1757 degrees of freedom
## (72 observations deleted due to missingness)
## Multiple R-squared:  0.009575, Adjusted R-squared:  0.008448
## F-statistic: 8.493 on 2 and 1757 DF, p-value: 0.0002134
```

```
summary(a8)
```

```
##
## Call:
## lm(formula = data$normalized_product_apgar ~ data$monpre + data$npvis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1897.44   -98.29   115.74   130.55   634.08
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1795.067     29.566  60.713 < 2e-16 ***
## data$monpre  -8.502       5.774  -1.472  0.14107
## data$npvis    6.313       1.936   3.261  0.00113 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 284.1 on 1757 degrees of freedom
## (72 observations deleted due to missingness)
## Multiple R-squared:  0.00981,    Adjusted R-squared:  0.008683
## F-statistic: 8.704 on 2 and 1757 DF,  p-value: 0.0001732
```

```
a9 = lm(data$normalized_combined_apgar~ data$npvis)
```

```
a10 =lm(data$normalized_product_apgar~ data$npvis)
```

```
summary(a9)
```

```
##
## Call:
## lm(formula = data$normalized_combined_apgar ~ data$npvis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -329.62  -10.01   13.80   15.62   68.72
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 394.7265     2.8167 140.136 < 2e-16 ***
## data$npvis   0.9108     0.2310   3.943 8.35e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 35.63 on 1759 degrees of freedom
## (71 observations deleted due to missingness)
## Multiple R-squared:  0.008763,    Adjusted R-squared:  0.008199
## F-statistic: 15.55 on 1 and 1759 DF,  p-value: 8.35e-05
```

```
summary(a10)
```

```
##
## Call:
## lm(formula = data$normalized_product_apgar ~ data$npvis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```



```
## -1899.19   -97.35   117.97   132.39   631.01
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1766.343     22.468   78.616 < 2e-16 ***
## data$npvis    7.211       1.842    3.914 9.43e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 284.2 on 1759 degrees of freedom
## (71 observations deleted due to missingness)
## Multiple R-squared:  0.008633, Adjusted R-squared:  0.00807
## F-statistic: 15.32 on 1 and 1759 DF, p-value: 9.429e-05
```

We did not see very good results with the APGAR score variations, so now let's try to normalize the baby's birth weight by APGAR.

```
data$weight_apgar_ratio = data$bwght/data$combined_apgarscores
a11 = lm(data$weight_apgar_ratio ~ data$npvis )
a12 = lm(data$weight_apgar_ratio ~ data$npvis + data$monpre)
summary(a11)
```

```
##
## Call:
## lm(formula = data$weight_apgar_ratio ~ data$npvis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -143.38  -22.18   -2.07   19.03  448.91
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 190.8039     3.0457  62.648 <2e-16 ***
## data$npvis    0.5288     0.2498   2.117  0.0344 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 38.53 on 1759 degrees of freedom
## (71 observations deleted due to missingness)
## Multiple R-squared:  0.002543, Adjusted R-squared:  0.001975
## F-statistic: 4.484 on 1 and 1759 DF, p-value: 0.03436
```

```
summary(a12)
```

```
##
## Call:
## lm(formula = data$weight_apgar_ratio ~ data$npvis + data$monpre)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -142.40  -22.50   -1.93   18.86  449.31
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 186.2113     4.0080  46.460 <2e-16 ***
## data$npvis    0.6709     0.2624   2.557  0.0107 *
## data$monpre    0.0000     0.0000   0.000  1.0000 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
## data$monpre    1.3835    0.7827    1.768    0.0773 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 38.51 on 1757 degrees of freedom
## (72 observations deleted due to missingness)
## Multiple R-squared:  0.004306,    Adjusted R-squared:  0.003173
## F-statistic: 3.799 on 2 and 1757 DF,  p-value: 0.02257
```

```
AIC(a11)
```

```
## [1] 17861.68
```

```
AIC(a12)
```

```
## [1] 17851.35
```

```
####these exclude NaN and Infs are not working :(
#data$weight_apgar_prod_ratio = data$bwght/data$product_apgarscores
#a13 = lm(data$weight_apgar_prod_ratio ~ data$npvis,na.action = NULL, singular.ok = #TRUE)
#a14 = lm(data$weight_apgar_prod_ratio ~ data$npvis + data$monpre, na.exclude(data), #singular.ok = TRUE)
#summary(a13)
#summary(a14)
#AIC(a13)
#AIC(a14)

#p <- ggplot(a14, aes(a14$residuals, a14$fitted)) + geom_point(size = 0.25) +
#  geom_smooth(method = "lm", se = FALSE) +
#  ggtitle("Resid vs Fitted")
```


Model 4: Problematic covariants

```
c1 = lm(data$lbwght ~ data$monpre + data$npvis + data$male +
        data$mwhite + data$fblck)
summary(c1)
```


```
##
## Call:
## lm(formula = data$lbwght ~ data$monpre + data$npvis + data$male +
##     data$mwhite + data$fblck)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.18137 -0.08572  0.02363  0.11920  0.45080
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7.959815   0.028796 276.418  < 2e-16 ***
## data$monpre  0.008995   0.004132   2.177   0.0296 *
## data$npvis   0.007604   0.001383   5.500 4.35e-08 ***
## data$male    0.021662   0.009702   2.233   0.0257 *
## data$mwhite  0.038981   0.019472   2.002   0.0454 *
## data$fblck   0.018740   0.026304   0.712   0.4763
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 0.2033 on 1757 degrees of freedom
## (69 observations deleted due to missingness)
## Multiple R-squared:  0.02158,    Adjusted R-squared:  0.01879
## F-statistic: 7.749 on 5 and 1757 DF,  p-value: 3.201e-07
```


```
plot(c1)
```




```
lm(data$lbwght ~ data$monpre + data$npvis + data$male + data$mwhite + data$f ..
Normal Q-Q
```



```
lm(data$lbwght ~ data$monpre + data$npvis + data$male + data$mwhite + data$f ..
```



lm(data\$lbwght ~ data\$monpre + data\$npvis + data\$male + data\$mwhite + data\$f ..
Residuals vs Leverage



lm(data\$lbwght ~ data\$monpre + data\$npvis + data\$male + data\$mwhite + data\$f ..

From all of the summaries, we can tell that the t-statistic for the `monpre` variable is not significant. Thus, we cannot trust this particular regressor, as the summary statistics on this regressor suggested.

Step 4: CLM and the Models

Step 5: Regression Tables and Model Analysis

Step 6: Causality

Biases and Limitation

This data is extremely biased in that no still births were included in our dataset. It is a sad fact in the United States that over 2 in 1,000 births are stillbirths. Since we do not know the prenatal care data for stillbirths, we cannot completely gauge how much prenatal care contributes to a child's health at birth.

In addition, it appears that there is little correlation between the Apgar score and the later health of the baby. The Apar is only meant to be used in the context of emergency situations. In this manner, looking at a baby's weight will give us deeper insight into the baby's overall health.

Step 7: Conclusion